


Archive



VLM caption
gemini-2.5-pro

Captions



“A black and white line drawing of the silhouette of a sports car.”



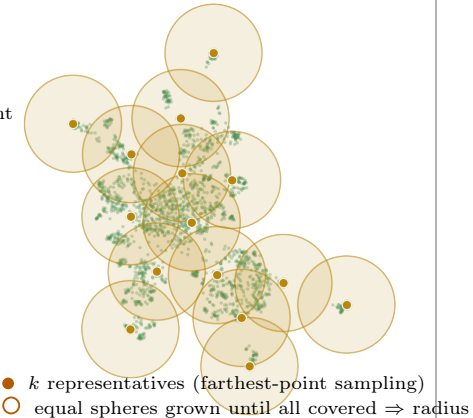
“A black and white drawing of a Stormtrooper helmet.”



“A minimalist black and white line drawing of a cat looking up.”

text
embedding
gemini-embedding

Caption-embedding space



Semantic Coverage = k -covering radius ($k = 100$) over caption embeddings
larger radius $\Rightarrow$ more semantically distinct forms